

Big Data

Assignment 2 Report

Author: Nazgul Salikhova (n.salikhova@innopolis.university)

Group: B22-AAI-02

Methodology

The assignment focuses on building a basic information retrieval pipeline using a distributed architecture. We leverage **Apache Spark** for data processing, **Hadoop HDFS** for storage, and **Apache Cassandra** for storing the indexed data and scores. The goal is to preprocess a document collection, extract relevant statistics, and compute BM25 scores for efficient document retrieval. The project is not done, because only 1st stage is ready - Indexing

Data Preparation

The documents are loaded into **HDFS**, where Spark jobs access them for tokenization, normalization, and stop-word removal. Each document is assigned a unique doc_id and its tokens are extracted for indexing.

1) Indexing

Using Spark, we compute the following:

- **Vocabulary statistics:** document frequency and total occurrences of each term.
- **Document statistics:** length, title, and average term frequency per document.
- **Inverted index:** mapping terms to the documents they appear in along with term frequency and word positions.
- **BM25 scores:** relevance scores for each (term, document) pair using the BM25 algorithm.

Cassandra Table Description

To store the information obtained during document indexing, a **keyspace** named search_index was created in Cassandra. It uses the NetworkTopologyStrategy replication strategy with one replica in datacenter1.

1. vocabulary

Stores statistics about individual terms across the document collection:

- term — the text value of the term (primary key).
 - document_frequency — the number of documents that contain the term.
 - total_occurrences — the total number of times the term appears in the entire collection.
2. **document_stats**
- Contains metadata for each document:
- doc_id — the unique identifier of the document (primary key).
 - title — the title of the document.
 - doc_length — the total number of words in the document.
 - avg_term_frequency — the average frequency of terms in the document (total term frequency divided by the number of unique terms).
3. **inverted_index**
- Represents a classic inverted index:
- term — the term (part of the composite primary key).
 - doc_id — the ID of the document where the term appears (second part of the composite key).
 - term_frequency — the number of times the term appears in the document.
 - positions — a list of word positions in the document where the term occurs.
4. **bm25_scores**
- Stores precomputed BM25 relevance scores for each term-document pair:
- term — the term text.
 - doc_id — the document ID.
 - score — the BM25 score representing the relevance of the document to the term.

These tables provide an efficient structure for implementing full-text search, term-based lookups using an inverted index, and document ranking based on relevance metrics.

How to Run

1. Clone the repository
2. Start the system using Docker Compose:
 - `docker-compose up -d`